

Chapter 3: Secure Links

While TCP provides a reliable way to send bytes over a network, the data is sent in the clear, making it vulnerable to interception by a middleman. To secure this communication, the **Transport Layer Security (TLS)** protocol is used. TLS runs on top of TCP, providing a secure channel that application layer protocols like HTTP can use.

In essence, TLS provides three core security guarantees:

- **Encryption:** Protects data from being read by unauthorized parties.
- **Authentication:** Verifies the identity of the communicating parties.
- **Integrity:** Ensures data has not been tampered with or corrupted.

3.1 Encryption

Encryption guarantees that data transmitted between a client and a server is obfuscated and can only be understood by them.

The process involves two types of cryptography:

1. Asymmetric Encryption:

- This method is used when the TLS connection is first established to negotiate a shared encryption secret.
- Each party generates a key pair (a private and a public key) and exchanges their public keys. The shared secret can be derived from this exchange without ever being sent over the wire.
- Asymmetric encryption is slow and computationally expensive, so it's only used for this initial key exchange.

2. Symmetric Encryption:

- After the shared secret is established, all subsequent communication is encrypted using symmetric encryption, which is fast and cheap.
- The shared key is periodically renegotiated to minimize the risk if a key were to be broken.

Note: Encrypting data has a CPU penalty, but this is negligible on modern processors with dedicated cryptographic instructions. Therefore, **TLS should be used for all communications**, including those not on the public internet.

3.2 Authentication

Authentication is the process of verifying that the server is who it claims to be. The server may also want to authenticate the client. TLS implements this using **digital signatures** based on asymmetric cryptography.

Certificates

The main challenge in authentication is for the client to trust that the public key shared by the server is authentic. This problem is solved using **certificates**.

- **Certificate:** A digital document that proves ownership of a public key. It includes:
 - Information about the owning entity.
 - An expiration date.
 - The public key itself.
 - A digital signature from the third-party entity that issued it.
- **Certificate Authority (CA):** The entity that issues a certificate.
- **Certificate Chain:** Certificates are linked together in a chain. A service's

certificate is signed by a CA, whose certificate might be signed by another CA, and so on. This chain ends with a **root CA**, which issues a self-signed certificate.

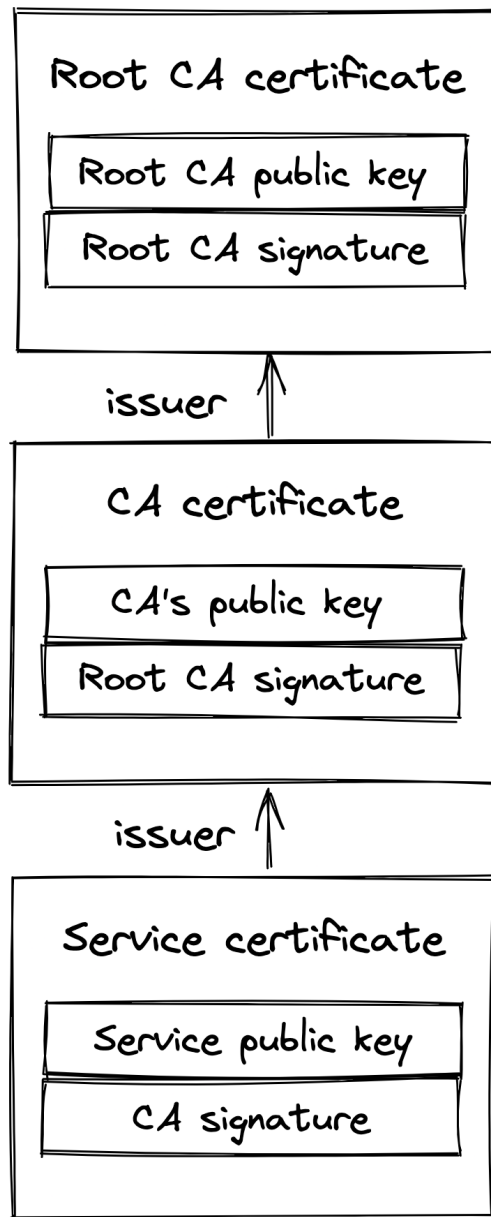


Figure 1: A certificate chain ends with a self-signed certificate issued by a root CA.

Trust Verification

- A client's device (e.g., operating system) maintains a **trusted store** containing certificates of trusted root CAs like *Let's Encrypt*.

- When a TLS connection is opened, the server sends its full certificate chain to the client.
- The client scans the chain until it finds a certificate issued by a CA that is in its trusted store.
- It then verifies each certificate in reverse order, checking the digital signature and expiration date until it reaches the server's own certificate. If all checks pass, the server is authenticated.

Important: A common and critical operational failure is letting a certificate expire. This will cause clients to fail connection attempts, potentially bringing down an entire application. It is crucial to have automation in place to monitor and auto-renew certificates.

3.3 Integrity

Integrity ensures that the data has not been altered or corrupted in transit. Even if data is encrypted, a middleman could still tamper with it (e.g., by swapping random bits).

- To protect against this, TLS calculates a **message authentication code (HMAC)** for each message using a secure hash function.
- When a process receives a message, it recomputes the HMAC and compares it to the one included in the message.
- If the codes do not match, it means the message has been tampered with or corrupted, and the message is dropped.

Note on TCP Checksums: While TCP uses a checksum for integrity, it is not 100% reliable and can fail to detect errors in rare cases (roughly 1 in 16 million to 10 billion packets). The TLS HMAC provides a much

stronger guarantee against data corruption.

3.4 Handshake

When a new TLS connection is established, the client and server perform a **handshake**. During this process, they:

1. **Agree on a Cipher Suite:** They negotiate the specific set of algorithms to use for:
 - Key exchange (for generating the shared secret).
 - Signatures (for signing certificates).
 - Symmetric encryption (for application data).
 - HMAC (for data integrity).
2. **Create a Shared Secret:** They use the agreed-upon key exchange algorithm to create the shared secret that will be used for symmetric encryption.
3. **Authenticate:** The client verifies the server's certificate. The server can optionally verify a client certificate as well.

This handshake process requires additional network round trips (typically 2 for TLS 1.2 and just 1 for TLS 1.3). This is another reason to place servers geographically close to clients and to reuse existing connections whenever possible to avoid this initial latency cost.