

# Chapter 6: System Models

To effectively reason about distributed systems, we use **system models**. A system model is a set of assumptions that define what can and cannot happen within a system. It simplifies reality by encoding expectations about three core components:

1. Communication links
2. Processes
3. Timing

These models allow us to analyze and verify the correctness of distributed algorithms without getting lost in the complexity of the underlying implementation technologies.

## Communication Link Models

These models define the assumed behavior of the network connections between processes.

- **Fair-loss Link:**
  - Assumes that messages may be lost or duplicated.
  - However, it guarantees that if a sender continuously retransmits a message, it will *eventually* be delivered.
- **Reliable Link:**
  - A stronger model that assumes a message is delivered *exactly once*, without any loss or duplication.
  - A reliable link can be implemented on top of a fair-loss link, for example, by using sequence numbers to de-duplicate messages at the receiver. TCP is an example of a protocol that implements a reliable link.

- **Authenticated Reliable Link:**

- Builds on the reliable link model.
- It makes the same assumptions but adds the guarantee that the receiver can authenticate the sender's identity. This can be implemented using a protocol like TLS on top of TCP.

## Process Failure Models

These models describe how processes are expected to behave, particularly concerning failures.

- **Arbitrary-fault Model (Byzantine Model):**

- The most extreme model, it assumes a process can deviate from its algorithm in any way.
- This includes crashing, sending incorrect data, or acting maliciously due to bugs or attacks.
- A key theoretical result is that a system can tolerate up to  $\frac{1}{3}$  of its processes being Byzantine faulty and still operate correctly.
- This model is typically used for safety-critical systems (e.g., airplane engines, nuclear power plants) and decentralized systems where no single entity has full control (e.g., Bitcoin).

- **Crash-recovery Model:**

- Assumes that a process follows its algorithm correctly but can crash at any time.
- After a crash, the process can restart, but it loses its in-memory state.

- **Crash-stop Model:**

- A simpler model that assumes a process follows its algorithm correctly but, if it crashes, it **never comes back online**.
- While unrealistic for software crashes, this model is useful for representing unrecoverable hardware faults and generally simplifies the design of algorithms.

## Timing Models

These models define the assumptions made about the time it takes for messages to be delivered and for processes to execute operations.

- **Synchronous Model:**

- Assumes that sending a message or executing an operation will never take more than a known, fixed amount of time.
- This model is **not very realistic** for most distributed systems, where network latency can be unpredictable and processes can be slowed by garbage collection cycles or page faults.

- **Asynchronous Model:**

- Assumes that sending a message or executing an operation can take an **unbounded** amount of time.
- While this model is simpler to reason about, many distributed computing problems are impossible to solve under this assumption, as an algorithm could get stuck waiting forever.

- **Partially Synchronous Model:**

- A realistic middle ground that assumes the system behaves synchronously *most of the time*, but allows for periods of asynchronous

behavior (e.g., high network latency).

- This model is generally considered representative enough of real-world systems.

## **The Book's Assumed Model**

For the most part, the algorithms presented in this book will assume a system model with:

- **Fair-loss links**
- **Crash-recovery processes**
- **Partial synchrony**

It is important to remember that all models are abstractions of reality. They are useful tools for reasoning, but they don't capture every real-world nuance. As you learn about algorithms, it is helpful to question the models' assumptions and consider how those algorithms might break in practice.