

Chapter 17

File Storage

1. Introduction to Managed File Stores

- **Problem:** Application servers have limited local disk space and can run out when storing many images, videos, and other static files.
- **Solution:** Utilize **managed file stores**, such as *AWS S3* or *Azure Blob Storage*, to store large static files.
- **Key Benefits:** These services are typically:
 - **Scalable**
 - **Highly available**
 - Offer **strong durability guarantees**
- **CDN Integration:** Files uploaded to a managed store can be configured for public URL access, enabling CDNs to point directly to them.
- **Outcome:** This approach allows for the *complete offloading* of storage and serving of static resources to managed services, significantly reducing the load on the application server.

2. Blob Storage Architecture (Focus on Azure Storage - AS)

Understanding the internal workings of distributed file stores is beneficial due to their crucial role in modern applications. This section focuses on Azure Storage (AS) as an example of a scalable cloud storage system with strong consistency, specifically its file abstraction (blob store).

A. Global Infrastructure

- AS consists of **storage clusters** distributed across *multiple regions world-wide*.
- Each storage cluster comprises multiple **racks of nodes**. Each rack is a separate unit with *redundant networking and power*.

B. Namespace and Access

- AS uses a **global namespace** based on domain names, composed of an *account name* and a *file name*.
- These two parts form a **unique URL** pointing to a specific file (e.g., `https://ACCOUNT_NAME.blob.core.windows.net/FILE_NAME`).
- The *account name* (customer-configured) is used by AS DNS to identify the specific **storage cluster** where the data resides.
- The *file name* is used by the identified cluster to locate the **node** responsible for that data.

C. Location Service (Global Control Plane)

- This central service manages:
 - Creation of new accounts.
 - Allocation of accounts to storage clusters.
 - Moving accounts between clusters for better *load distribution*.
- **New Account Creation Process:**
 1. The location service selects a suitable cluster for the new account (in a specific region requested by the customer) based on current *load information*.
 2. It updates the chosen cluster's configuration to begin accepting requests

for this new account.

3. It creates a new *DNS record* that maps the account name to the cluster's public IP address.

D. Storage Cluster Architecture (Three Layers)

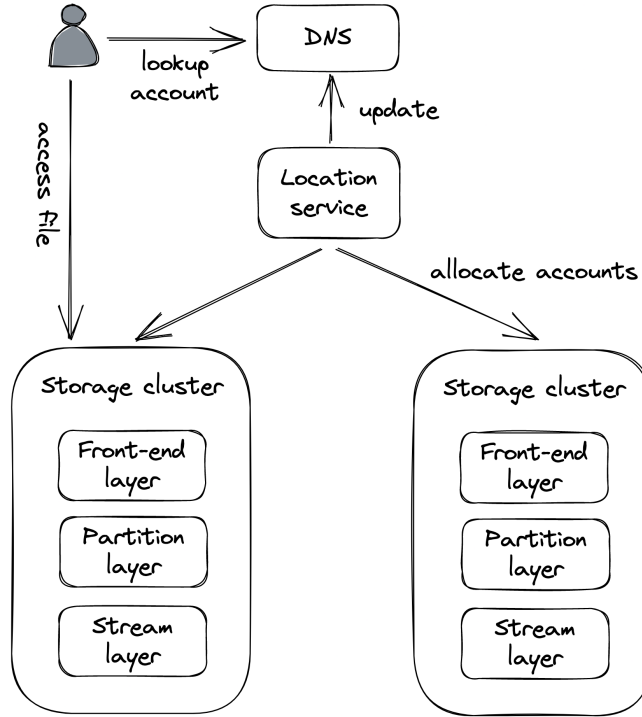


Figure 1: A high-level view of Azure Storage's architecture

A storage cluster in AS is architecturally composed of three distinct layers:

1. Stream Layer:

- Implements a *distributed append-only file system*. Data is stored in structures called **streams**.
- Internally, a stream is a sequence of **extents**. The *extent* is the fundamental unit of replication.
- Writes to extents are replicated *synchronously* using **chain replication**.
- **Stream Manager (Control Plane):**




Figure 2: The stream layer uses chain replication to replicate extents across storage servers

- Responsible for assigning an extent to a specific *chain of storage servers* within the cluster.
- When allocating a new extent, it provides the client (partition layer) with the list of storage servers in the chain for that extent.
- The client caches this information and uses it to send future writes directly to the *primary server* in the chain.
- Handles *unavailable or faulty extent replicas* by creating new ones and reconfiguring their replication chains.

2. Partition Layer:

- This layer translates *high-level file operations* (like read, write, delete on a file) into *low-level stream operations* (on extents).
- **Partition Manager (Control Plane):**
 - Manages a *large distributed index* of all files stored in the cluster.
 - Each index entry contains metadata (e.g., account and file name) and a pointer to the actual data in the stream service (which includes a list of extents, plus offset and length within those extents).
 - It *range-partitions* this index and maps each index partition to a




Figure 3: The partition manager range-partitions files across partition servers and rebalances the partitions when necessary

specific **partition server**.

- Responsible for *load-balancing* these index partitions across partition servers, including *splitting* partitions when they become too “hot” (frequently accessed) and *merging* “cold” ones.
- This layer also *asynchronously replicates accounts across different clusters* in the background. This is used for migrating accounts for load balancing and for disaster recovery purposes.

3. Front-End Layer:

- A *stateless service* that acts as a **reverse proxy**.
- It *authenticates incoming requests*.
- It routes these requests to the *appropriate partition server* by consulting the mapping managed by the partition manager.

E. Architectural Significance & Consistency

- The architecture of AS demonstrates the application of various *scalability patterns* in a large-scale, concrete system.
- **Historical Note on Consistency:** Azure Storage was designed and built with *strong consistency* from its inception. In contrast, AWS S3 initially offered eventual consistency for some operations and started providing strong read-after-write consistency in 2021.