

Chapter 19

Data Storage

1. Introduction: The Database Bottleneck

- Even with stateless application servers scaled out behind a load balancer, the **relational database**, often hosted on a *single machine*, can become the next performance bottleneck as application load and requests to the database increase.

2. Database Replication (Scaling Reads & Increasing Availability)

- **Purpose:** To increase the *read capacity* and *availability* of the database.
- **Common Topology: Leader-Follower**
 - **Writes (Updates, Inserts, Deletes):** Clients send writes *exclusively to the leader* replica.
 - The leader persists these changes to its **write-ahead log (WAL)**.
 - **Followers (Replicas):** Connect to the leader and *stream log entries* from it, applying (committing) them locally.
 - Log entries have *sequence numbers*, allowing followers to disconnect/reconnect and resume replication from where they left off by communicating the last processed sequence number to the leader.
- **Benefits of Replication:**
 - **Increased Read Capacity:** *Read-only followers* can be placed behind




Figure 1: Single leader replication

a load balancer to distribute read queries.

- **Increased Availability:**

- * A load balancer can automatically remove a *faulty replica* from the pool if it's unhealthy or unavailable.
- * If the *leader fails*, a replica can be reconfigured (promoted) to take its place.

- **Workload Isolation:** Individual followers can be used for specific, potentially intensive workloads (e.g., periodic analytics queries) without impacting the performance of the leader or other replicas serving live traffic.

- **Replication Modes:**

- **Fully Asynchronous Replication:**

- * The leader receives a write, broadcasts it to followers, and *immedi-*

ately sends a response back to the client without waiting for follower acknowledgments.

- * **Pros:** Minimizes client response time.

- * **Cons:** *Not fault-tolerant*; if the leader crashes after acknowledging a write but before broadcasting it, data loss can occur.

– **Fully Synchronous Replication:**

- * The leader *waits for the write to be acknowledged by all followers* before returning a response to the client.

- * **Pros:** Stronger data consistency, no data loss on leader failure if followers are up-to-date.

- * **Cons:**

- *Performance Cost:* A single slow replica increases the response time for every write request.

- *Availability Risk:* If any replica is unreachable, the data store becomes unavailable for writes.

- *Not Scalable:* The more followers, the higher the probability of at least one being slow or unavailable.

– **Combined (Semi-Synchronous) Replication:**

- * Common in practice. Relational databases like *PostgreSQL* often support configuring individual followers for synchronous replication (default is usually asynchronous).

- * **Example Use Case:** Have a single *synchronous follower* designated as an up-to-date backup of the leader.

- * **Benefit:** If the primary leader fails, failover can occur to the synchronous follower *without any data loss*.

- **Failover Mechanism (Conceptual Steps):**

1. Detect when the leader has failed.
 2. Promote the synchronous follower to be the new leader.
 3. Reconfigure other replicas to follow the newly promoted leader.
 4. Ensure client requests are routed to the new leader.
- **Managed Solutions:** Cloud database services like *AWS RDS* or *Azure SQL Database* often provide features like read replicas and automated failover, along with automated patching and backups, out of the box.
 - **Limitations of Replication:**
 - Primarily helps scale out *reads*, not writes (writes still go to the single leader).
 - The *entire database must still fit* on the leader's single machine. While moving some tables to other databases can delay this issue, it's not a fundamental solution for very large datasets.

3. Partitioning (Scaling Reads, Writes & Handling Large Datasets)

- **Purpose:** Overcomes the write scaling and single-machine size limitations of simple replication. Allows scaling out a database for *both reads and writes*.
- **Application-Layer Partitioning (for Traditional RDBMS):**
 - Traditional (centralized) relational databases *generally don't support partitioning out of the box*.
 - Implementing it at the application layer is, in principle, possible but *highly challenging* and adds significant complexity:
 - * **Data Distribution:** Deciding *how to partition* the data among the database instances.

- * **Rebalancing:** Moving data when a partition becomes too hot (frequently accessed) or too large.
- * **Cross-Partition Queries:** Queries spanning multiple partitions (e.g., aggregations, joins) need to be *split into sub-queries*, and their responses must be combined.
- * **Distributed Transactions:** Supporting atomic transactions across partitions requires implementing a distributed transaction protocol (e.g., *Two-Phase Commit (2PC)*).
- * **Combined Complexity:** Adding replication to an application-layer partitioning scheme further increases the difficulty.
- **The Fundamental Problem with Traditional Relational Databases:**
 - They were designed assuming they would fit on a *single, powerful (“beefy”) machine*.
 - Consequently, they support features that are *hard to scale* in a distributed manner, such as ACID transactions and complex joins across the entire dataset.
 - **Historical Context:** Designed when disk space was expensive, prioritizing *data normalization* to reduce storage footprint, even if it meant costly denormalization via joins at query time.
 - **Modern Context:** Storage is now relatively cheap, while CPU time is not.
- **Benefits of Normalization (Beyond Storage):**
 - **Data Integrity:** If data is duplicated, updates must occur in all places. With normalized data, an update needs to happen in only one place.

4. NoSQL Databases

- **Origins:** Emerged in the early 2000s from large tech companies building bespoke data storage solutions designed from the ground up for *high availability and scalability* (inspired by papers like Google’s Bigtable and Amazon’s Dynamo).
- **Initial Naming:** Termed “NoSQL” because the first generation often didn’t support SQL. This designation is now *misleading*, as many NoSQL stores have evolved to support SQL-like query languages or other rich features.
- **Key Differences from Traditional Relational Databases (RDBMS):**
 - **Consistency Models:** NoSQL stores often embrace *relaxed consistency models* (e.g., eventual consistency, causal consistency) to achieve high availability, whereas RDBMS typically support stronger models like strict serializability.
 - **Joins:** Generally *do not provide join operations*.
 - **Data Model:** Rely on *unnormalized data*, often represented as:
 - * **Key-Value Pairs:** An opaque key maps to an opaque value (sequence of bytes).
 - * **Documents (e.g., JSON):** A key maps to a (possibly hierarchical) document, often without a strictly enforced schema. Unlike pure key-value stores, documents are *interpreted and indexed*, allowing queries based on their internal structure.
 - **Transactions:** Due to native support for partitioning for scalability, NoSQL stores have *limited support for transactions*. For example, Azure Cosmos DB currently supports transactions scoped only to *individual partitions*. The need for complex transactions is reduced when data is

denormalized.

- **Using NoSQL for Relational Data:**

- While NoSQL stores can model relational data, trying to use them *as if they were traditional relational databases* will likely result in the “worst of both worlds.”
- When used correctly, NoSQL can handle many use cases of traditional RDBMS while being *essentially scalable from day one*.

- **Core Requirement for Efficient NoSQL Use:**

- **Know the access patterns upfront** and model the data accordingly. This is crucial.
- **Amazon DynamoDB Example:**
 - * **Abstraction:** A table containing items. Each item can have different attributes but *must* have a primary key.
 - * **Primary Key:** Can be a single attribute (the *partition key*) or two attributes (the *partition key* and the *sort key*).
 - **Partition Key:** Dictates how data is partitioned and distributed across nodes.
 - **Sort Key:** Defines how data is sorted *within a partition*, enabling efficient range queries.
 - * **Replication:** DynamoDB creates *three replicas* for each partition and uses state machine replication to keep them synchronized. Writes are acknowledged to the client when *two out of three* replicas have received the write.

* **Read Consistency:** Reads can be either *eventually consistent* (from any replica) or *strongly consistent* (querying the leader replica).

* **API Support:**

- CRUD (Create, Read, Update, Delete) operations on single items.
- Querying multiple items that share the *same partition key* (optionally specifying conditions on the sort key).
- Scanning the entire table (less efficient, for bulk operations).

* **No Joins by Design:** Joins don't scale well. Data should be modeled to avoid needing them.

* **Modeling Example (Customers and Orders):**

- Suppose the most common access pattern is retrieving orders for a customer, sorted by date. The table might have **customer ID** as the partition key and **order creation date** as the sort key.

| Partition Key | Sort Key | Attribute | Attribute |
|---------------|------------|---------------|-----------------|
| jonsnow | 2021-07-13 | OrderID: 1452 | Status: Shipped |
| aryastark | 2021-07-20 | OrderID: 5252 | Status: Placed |
| branstark | 2021-07-22 | OrderID: 5260 | Status: Placed |

- If customer full name is also needed, entities of multiple types (customers, orders) can be stored in the same table:

| Partition Key | Sort Key | Attribute | Attribute |
|---------------|------------|---------------|-----------------|
| jonsnow | 2021-07-13 | OrderID: 1452 | Status: Shipped |

| Partition Key | Sort Key | Attribute | Attribute |
|---------------|------------|----------------------|----------------|
| jonsnow | jonsnow | FullName: Jon Snow | Address: ... |
| aryastark | 2021-07-20 | OrderID: 5252 | Status: Placed |
| aryastark | aryastark | FullName: Arya Stark | Address: ... |

- Because a customer and its orders have the same partition key, a single query can retrieve all related entities. This structures the table based on access patterns to avoid joins.

* **Secondary Indexes:** DynamoDB supports secondary indexes for more complex access patterns:

- *Local Secondary Indexes (LSIs):* Allow alternate sort keys on the same table (within the same partition key).
- *Global Secondary Indexes (GSIs):* Allow different partition and sort keys, effectively creating a new “view” of the data. Index updates for GSIs are asynchronous and *eventually consistent*.

- **NoSQL Flexibility Misconception:**

- It’s a common misconception that NoSQL data stores are more flexible because they scale seamlessly without upfront data modeling.
- The reality is that NoSQL *requires significantly more attention* to how data is modeled based on access patterns.
- Because they are tightly coupled to access patterns, NoSQL stores can be *less flexible* than relational databases if those patterns change.

- **Key Takeaway:** Using a NoSQL data store effectively *requires identifying the access patterns upfront to model the data accordingly*.

5. NewSQL Databases

- **Latest Trend:** Aim to combine the *scalability of NoSQL* with the *ACID guarantees* often associated with relational databases.
- **Consistency vs. Availability:** While many NoSQL stores prioritize availability over consistency in the face of network partitions (CAP theorem trade-off), NewSQL stores often *prefer consistency*.
- **Argument for NewSQL:** Proponents argue that for many applications, the reduction in availability caused by enforcing strong consistency is hardly noticeable, especially since perfect 100% availability is practically impossible anyway (availability is defined in “nines”).
- **Examples:** Well-known NewSQL data stores include *CockroachDB* and *Spanner*.