

Chapter 16: Partitioning

1. What is Partitioning and Why is it Needed?

- **Definition:** When an application's data volume grows too large to fit on a single machine, it needs to be split into smaller pieces called **partitions** or **shards**. Each partition is small enough to fit on an individual node.
- **Primary Goal:** To manage *large datasets* that exceed the capacity of a single machine.
- **Additional Benefit:** Increases the system's capacity for handling requests because the *load of accessing data is spread* across multiple nodes.

2. The Role of a Gateway Service

- When a client sends a request to a partitioned system, it needs to be routed to the *correct node(s)*.
- A **gateway service** (like a reverse proxy) is typically responsible for this routing, knowing how the data is mapped to partitions and nodes.
- This data-to-partition mapping is usually maintained by a *fault-tolerant coordination service* (e.g., etcd or Zookeeper).

3. Complexities and Drawbacks of Partitioning

Partitioning is not without its challenges and introduces significant complexity:

- **Gateway Requirement:** A gateway service is *necessary* to direct




Figure 1: A partitioned application with a gateway that routes requests to partitions (Figure 16.1)

requests to the appropriate nodes.

- **Data Aggregation:** To roll up data across different partitions (e.g., for a “group by” operation), data must be fetched from multiple partitions and then aggregated, which *adds complexity*.
- **Cross-Partition Transactions:** Transactions that need to atomically update data spanning multiple partitions *limit scalability*.
- **Hotspots:** If a partition is accessed much more frequently than others, the system’s ability to scale is limited.
- **Dynamic Resizing (Rebalancing):** Adding or removing partitions at runtime is *challenging* because it requires moving data across nodes.

4. Partitioning and Caches

- Caches are *well-suited* for partitioning because they avoid many of the common complexities.
- For instance, caches generally *don’t require atomic updates* across parti-

tions or complex aggregations spanning multiple partitions.

5. Prerequisite for Key Partitioning

- A fundamental requirement for partitioning key-value data is that the number of possible keys must be *very large*.
- Keys with a small set of possible values (e.g., a boolean key with only two values) are *not suitable* for partitioning as they allow for at most two partitions.

6. Methods of Mapping Key-Value Data to Partitions

There are two primary ways to map key-value pairs to partitions: range partitioning and hash partitioning.

6.1. Range Partitioning

- **Definition:** Splits data by *key range* into lexicographically sorted partitions (as shown in Figure 16.2).
- **Performance:** To make range scans fast, each partition is generally stored in sorted order on disk.
- **Challenges:**
 - **Picking Boundaries:**
 - * Evenly splitting the key range works well if key distribution is *uniform*.
 - * If not uniform (like words in a dictionary), partitions can become *unbalanced*, with some having significantly more entries




Figure 2: A range-partitioned dataset (Figure 16.2)

than others.

- **Hotspots:** Certain access patterns can lead to hotspots. For example, if data is range-partitioned by date, all requests for the current day might hit a *single node*.
 - * Workaround: Adding a random prefix to the partition keys can help, but it increases complexity.
- **Rebalancing (Adding/Removing Nodes):**
 - **Need:** When data size or request volume changes, nodes need to be added or removed to balance the load and manage costs. This process is called **rebalancing**.
 - **Goal:** Rebalancing should *minimize system disruption* and the amount of data transferred, as the system needs to continue serving requests.
 - **Static Partitioning:**
 - * Create *many more partitions* than initially needed and assign multiple partitions to each node. The number of partitions remains fixed over time.
 - * When a new node is added, some partitions are moved from

existing nodes to the new one to maintain balance.

* **Drawbacks:**

- The number of partitions is *fixed* and hard to change.
- Getting the initial number of partitions right is difficult: too many can add overhead and decrease performance; too few can limit scalability.
- Some partitions might still become hotspots if accessed much more than others.

– **Dynamic Partitioning:**

- * Partitions are created *on demand*.
- * The system starts with a single partition. When it grows too large or becomes too hot, it's split into two sub-partitions (approximately half the data each), and one sub-partition is moved to a new node.
- * Conversely, if two adjacent partitions become small or “cold” enough, they can be merged.

6.2. Hash Partitioning

- **Definition:** Uses a *hash function* to deterministically map a key to a seemingly random number (a hash) within a defined range (e.g., 0 to $2^{64} - 1$). This ensures keys' hashes are distributed *uniformly* across the range. A subset of these hashes is then assigned to each partition (as shown in Figure 16.3).
 - Example: $\text{hash}(\text{key}) \bmod N$, where N is the number of partitions.
- **Benefit:** Generally ensures partitions contain a *relatively similar number*




Figure 3: A hash-partitioned dataset (Figure 16.3)

of entries.

- **Challenges:**

- **Hotspots:** Does *not eliminate hotspots* if the access pattern is non-uniform. If a single key is accessed very frequently, the node hosting its partition can become overloaded.

- * Solutions: Further split the hot partition (increasing N), or split the hot key into sub-keys (e.g., by prepending a random prefix).

- **Rebalancing with Modulo Operator:** When a new partition is added using the $\text{hash}(\text{key}) \bmod N$ approach, *most keys have to be moved* (shuffled) to different partitions because their assignment changes. This shuffling is very expensive due to network bandwidth and resource consumption.

- * Ideally, adding a partition should only require shuffling K/N keys (where K is total keys, N is number of partitions).

- **Consistent Hashing:**

- A widely used hashing strategy that *minimizes data shuffling* during rebalancing.

- **How it works:** A hash function randomly maps both partition identifiers and keys onto an *imaginary circle*. Each key is assigned to the *closest partition* that appears on the circle in a clockwise direction (see Figure 16.4).




Figure 4: With consistent hashing, partition identifiers and keys are randomly distributed around a circle, and each key is assigned to the next partition that appears on the circle in clockwise order (Figure 16.4)

- **Adding a new partition:** When a new partition is added to the circle, *only the keys that now map to this new partition* (due to proximity on the circle) need to be reassigned (as shown in Figure 16.5).
- **Main Drawback of Hash Partitioning (vs. Range Partitioning):**
 - The *sort order of keys across partitions is lost*, which is required to efficiently scan all the data in order.
 - However, data *within* an individual partition can still be sorted based on a secondary key.




Figure 5: After partition P4 is added, the key ‘for’ is reassigned to P4, but the assignment of the other keys doesn’t change (Figure 16.5)